

Data Dictionary

Content Sunsetting Programme Tracker · Technical documentation · Ndivhuwo
ndivhuwojse@gmail.com · linkedin.com/in/ndivhuwo-makhavhu · github.com/MainDevWork

What the Content Sunsetting Programme Tracker records, and why it records it that way.

Start here

A large organisation's website grows one page at a time, usually over many years and always faster than anyone retires anything. Eventually it holds far more pages than any single person can account for. That is normally harmless, until the site has to move to a new content management system, at which point somebody has to decide, page by page, what gets carried forward and what gets switched off.

The blocker is almost never technical. It is that **nobody can say who is responsible for a given page**, so nobody can say whether it is safe to delete. Ask around and a page turns out to belong to a team that reorganised, or to someone who left, or to no one at all. Without an owner there is no one to ask, no one to sign off, and no defensible record of why a page was removed.

This tracker exists to remove that blocker. It makes **every URL carry a named owner, a decision, a reason and a date**, in one model, so the question "can we delete this?" has an answer that someone is accountable for.

The words this document uses

Four terms are used throughout in a specific sense:

- **Sunsetting**: retiring a web page, taking it out of service rather than editing it. The word is used in preference to "deleting" because several of the available outcomes are not deletions.
- **Custodian**: the named person accountable for a page. Not its author and not its approver: the person who answers for whether it should still exist.
- **Sitemap**: the machine-readable index a site publishes to tell search engines and analytics tools which pages exist. A page missing from it is reachable if you know the address, but is largely invisible, which makes it both a strong retirement candidate and a page with no traffic evidence.
- **Media Library**: handles by which stored files (documents, images) are reachable in the content management system. They are URLs, but they are *not web pages*, and they outnumber real pages heavily. Keeping them separate from any headline count is the single most important distinction in this model.

What the tracker is made of

Three linked record sets and a dashboard:


Record set	The question it answers
Webpage Custodianship Records	Who is responsible for this page?
Sunsetting Outcome Records	What was decided for it, why, and when?
URL Performance Dataset	What is it worth in traffic?

A two-page **Power BI dashboard** reads all three. A separate **Excel response form** is issued to each division so it can state what it wants done with its own pages; those responses are translated into the tracker.

Everything joins on one key: [URL Links](#), the page's web address. That single decision is what makes the estate legible: ownership, decision and traffic for any page are one lookup apart.

About the data. The schema, the controlled vocabularies, the join logic and the measures documented here are real. **Every value shown is synthetic**, generated from a fixed seed against a fictitious exchange, Sableridge Exchange (SRX), www.sableridge.co.za. Names, URLs, divisions, decisions and traffic figures are invented. No figure in this document describes any real organisation's website, staff or decisions.

1. The data model



URL Links is the only key in the model. It holds the absolute page URL, stored identically in all three files.

Cardinality. Custodianship and outcomes are strictly one-to-one: the same URLs, unique in both. Performance is a proper subset: most URLs carry a traffic record and some do not.

Why the subset exists, and why it constrains the analysis. A URL absent from the sitemap is not reachable by the analytics crawl, so it has no performance row. The same is true of pages published after the measurement window closed.

That has a consequence worth stating plainly, because it is easy to get wrong: **the pages most likely to be retired are precisely the pages with no traffic evidence.** **Very Low Page Views** is a sanctioned reason for decommissioning, but it cannot be evidenced from the traffic table for a page that is missing from it. For that population, *absence of a performance row is itself the signal*, not a low number in one. Reading the traffic table as though a missing row meant zero views would quietly invert the finding.

2. Webpage Custodianship Records

Webpage_Custodianship_Records.xlsx , sheet Webpage_Demographic , plus a Lists sheet holding the drop-down sources.

This is the ownership register, and it is where the custodianship verification process writes its results back.

#	Field	Type	Required	Role
1	WCR_ID	Text	Yes	Surrogate key, WCR<n> . Row identity only, never join on it.
2	URL Links	Text (absolute URL)	Yes	Join key. Unique across the file.
3	Description	Text	Yes	Human-readable page label. Deliberately not unique.
4	URL_Validation_Flag	Controlled, 3 values	Yes	Sitemap presence at audit. Governs whether a traffic row can exist.
5	Business Unit	Controlled	Yes	Owning team. Rolls up into Business Divisions .
6	Website Section	Controlled, 4 values	Yes	Which part of the estate the URL sits in. Slicer on both dashboard pages.
7	Business Divisions	Controlled	Yes	Owning division. Drives the division bar chart and slicer.
8	Webpage Custodian	Text	Yes	Named owner, or a system label.
9	Custodian Confirmed	Yes / No	Yes	Whether the named owner has accepted ownership.

Description is a display field, not an identifier. Many URLs share a description, because a description names the *kind* of page rather than the instance: every asset filed under one source page carries that page's label. Grouping or joining on it will silently collapse rows. URL Links is the identifier; Description exists so a human reading a row knows what they are looking at.

Webpage Custodian holds people and system labels in the same column. Most values are named individuals. Three are not people at all: they are how the model records content that no individual owns:

Label	What it means
Media Library	A stored-file handle, administered centrally and never authored by a division.
Digital Team	Site infrastructure, or a page no division claims. The decision rests with the digital team.
Yet to be Confirmed	Ownership genuinely unknown. Always paired with Custodian Confirmed = No .

This is deliberate. Every URL must resolve to an accountable party, and "nobody" is not an available answer. An unowned page with no marker is exactly the page that gets missed. The three labels make unowned content *visible and countable* rather than absent. The cost is that a distinct count over this column is a count of accountable parties, not a headcount; that is recorded with the measures in §7.

The same pattern repeats in the unit and division columns. Media Library and Yet to be Confirmed also appear as a business unit and a division. They are not parts of any organisation: they are how the model records content sitting outside the divisional structure entirely. They appear on the Lists sheet like any other permitted value, because the data legitimately contains them and a validation list that omits a value the data contains is worse than no validation list.

Webpage Custodian deliberately has no drop-down. It is a staff roster, which changes with joiners and leavers, not a controlled vocabulary. Validating it against a fixed list would block a legitimate new owner from being recorded, and

being unable to record the truth is a worse failure than a typo. Controlled lists are used where the vocabulary is a closed set by design; they are not used where the set is genuinely open.

Organisational vocabularies here are synthetic. The divisions, business units and custodian names in the published files are invented, and the unit-to-division mapping is a plausible fabrication. What is real is the *shape*: a controlled two-level hierarchy in which every unit resolves to exactly one division, so a division roll-up returns the same answer whether it is computed through `Business Unit` or read straight off `Business Divisions`. Rule 9 in Appendix A exists to keep it that way.

3. Sunsetting Outcome Records

Sunsetting_Outcome_Records.xlsx, sheet Sheet1, plus a Lists sheet.

This is the audit trail: what was decided for each page, why, and when.

#	Field	Type	Required	Role
1	SOR_ID	Text	Yes	Surrogate key, SOR<n>.
2	URL Links	Text (absolute URL)	Yes	Join key. Unique.
3	Description	Text	Yes	Display label, mirrored from custodianship.
4	Page Created	Date	No	Publication date. Reserved; see below.
5	Webpage Custodian	Text	Yes	Denormalised copy of the owner.
6	Date of Sunsetting Discussion	Date	No	When the decision was taken. Presence is meaningful.
7	Sunsetting Procedure	Controlled, 10 values	Yes	The decision. Validated drop-down.
8	Reason	Controlled, 5 values	No	Justification. Conditional; see §5.
9	Comments	Free text	No	Narrative reasoning behind the decision.

Date of Sunsetting Discussion is a status field in disguise. A blank here does not mean "date unknown". It means *the discussion has not happened yet*. Every URL still sitting at **Yet To Be Confirmed** is blank here by construction. This is why the field is optional and why nothing back-fills it: writing a plausible date into a page that was never discussed would corrupt the one measure that tracks how far the programme has actually got.

Page Created is defined and reserved. Page age is a strong independent signal for retirement, and the field is in the schema so that signal has somewhere to live. Populating it from content-management-system metadata is the highest-value extension available to this model: it would let age be used alongside traffic and ownership rather than in place of them, and it is the natural next iteration.

Webpage Custodian is denormalised on purpose. It appears in both record sets rather than being normalised away, because the two files are maintained by different people on different cycles (custodianship by the verification process, outcomes by the discussion schedule), and each has to be readable standing alone. Normalising would be tidier and would make the outcome file unusable without a join, which is the wrong trade for an instrument that people open in Excel. The accepted cost is that two copies can drift; the custodianship file is authoritative, and a reconciliation check between them is a sensible addition.

Comments carries the narrative reasoning behind a decision: the sentence explaining why a page was kept or dropped when the controlled **Reason** does not capture it. Free-text reasoning is the one field that cannot be meaningfully generalised, so the published values are generic stand-ins drawn from a small template pool per procedure. They demonstrate the field's purpose and typical length and nothing else.

4. URL Performance Dataset

URL_Performance_Dataset.xlsx , sheet Valid_URL_Performance .

What each reachable page is worth in traffic. This is the evidence base behind the Very Low Page Views reason, and it is the only one of the three files sourced from an external system rather than from the programme itself.

#	Field	Type	Required	Role
1	URL Links	Text (absolute URL)	Yes	Join key. Unique. Every value also exists in custodianship.
2	Description	Text	Yes	Display label, mirrored from custodianship.
3	Page Views	Integer	Yes	Views over the measurement window.
4	Average Session Duration (s)	Decimal	Yes	Mean seconds on the page.
5	Bounce Rate	Decimal, 0 to 1	Yes	Share of sessions ending on this page without further navigation.

This table covers only pages the analytics crawl can reach, which is why it is smaller than the other two and why \$1's warning matters: a URL's absence from this file is information, not a gap.

Traffic is power-law distributed: one dominant landing page and a very long tail, with a large share of URLs attracting almost no views at all. That is a property of the web rather than of any particular organisation, and the synthetic data reproduces it faithfully because Very Low Page Views is only a usable rationale when measured against a realistic distribution. Against a flat distribution it would either catch everything or nothing.

5. Controlled vocabularies

These are the real vocabularies, and they are the reason the dashboard can roll anything up at all. Decision values are constrained to a fixed list rather than typed freely, because a free-text decision column produces a chart with one bar per spelling: a misspelled value is invisible until somebody types it correctly, at which point the chart silently gains a second category and every percentage on the page moves. Constraining the list makes that class of error impossible rather than detectable.

Sunsetting Procedure: 10 values

The decision taken for a URL. Validated against the `Lists` sheet; a value outside this list cannot be entered.

Value	Meaning	Terminal?
To Be Decommissioned	Page will be retired and not carried forward.	Yes
To Be Migrated	Active page, carried to the new system as-is.	Yes
To Be Migrated (Content Update)	Carried forward, but the content is refreshed first.	Yes
Move To Group Site	Corporate rather than operational content; belongs on the group site.	Yes
Move To Client Portal	Relevant only to authenticated clients; belongs behind login.	Yes
Move To Clearing House Site	Clearing-related content; belongs on the clearing house site.	Yes
Temporarily Unpublished	Withdrawn pending a decision; restorable if a use case appears.	No
Another Discussion to be Held	Reviewed, decision deferred to a further session.	No
Digital Team Discretion	No division claims it; the decision rests with the digital team.	No
Yet To Be Confirmed	Default state. No decision has been taken.	No

Terminal means the decision is final and the page can be actioned.

The four non-terminal states are the point of the vocabulary, not padding. A simple confirmed/unconfirmed field would collapse "we discussed this and deferred it", "no division claims it" and "nobody has looked at it yet" into a single bucket. Those are three different problems needing three different interventions (a follow-up session, an escalation, and scheduling respectively), and a vocabulary that cannot tell them apart cannot be used to manage the work. The three `Move To ...` values exist for the same reason: "not staying here" and "being deleted" are routinely conflated, and they have opposite consequences for the content.

The `Move To ...` destinations are generalised in this published copy; the routing logic is real, the destination names are not disclosed.

Reason: 5 values

Value	Applies to
Page Info Irrelevant	To Be Decommissioned , Temporarily Unpublished
No Use Case	To Be Decommissioned
Very Low Page Views	To Be Decommissioned , Temporarily Unpublished
Custodian Request Decommission	To Be Decommissioned

Value	Applies to
No Reason / To Be Migrated	All four migrate/move procedures

Reason is conditional by design, and the pairings are enforced in the data rather than left to whoever is typing:

- **Always blank** for Another Discussion to be Held , Digital Team Discretion and Yet To Be Confirmed . No decision has been taken, so there is nothing to justify, and an empty justification against a non-decision is the correct state rather than a missing value.
- **Constrained elsewhere.** A page marked for migration is not decommissioned for having no use case, so the model will not record that pairing.

No Reason / To Be Migrated looks like a null and deliberately is not one. It states explicitly that no justification is required here, which keeps a genuine blank meaning "justification missing", a distinction that matters the moment the field is used to audit decision quality rather than just to display it.

URL Validation Flag: 3 values

Value	Meaning
In Sitemap	Present in the sitemap at audit. Reachable, and carries a traffic record.
NOT in Sitemap - Potential Sunset	Reachable by URL but unlisted. A standing retirement candidate.
In Sitemap - New (not in original audit)	Published after the original audit. May predate the traffic window.

Website Section: 4 values

Business · Instruments · Announcements · Media Library

This field carries the model's most important distinction, restated here because it changes every headline number:

Media Library URLs are stored-file handles, not web pages. They dominate any mature content estate by count, so a raw URL total reported without separating them overstates the size of the public site several times over. Every headline count on the dashboard is sliceable by this field for exactly that reason, and no figure should be quoted without saying which side of this line it falls on.

Decision: 3 values

Used only on the divisional response form, Custodian_Page_URLs_Template.xlsx : Keep · Consolidate · Sunset .

This vocabulary is deliberately coarser than the tracker's ten procedures. A division is being asked what it wants for its content, not to select a programme procedure. Asking a respondent to choose between Move To Group Site and Digital Team Discretion would be asking them to make a decision that is not theirs and that they lack the context to make. Translating a coarse answer into a precise procedure is a digital-team judgment, and treating it as an explicit, documented step is what keeps that judgment visible in the audit trail rather than buried in whoever happened to key the response.

6. How the instrument is designed

Five principles govern the collection instruments and the model. Each exists to prevent a specific, well-known failure mode.

One question, one answer. The response form asks for a decision in a single validated `Decision` column, not as a set of independent tick-boxes. Independent booleans let a respondent tick two contradictory answers, or none, and make "not answered" indistinguishable from "answered no", which is the difference between a response you can chase and a response you have already miscounted. One constrained column makes exactly one answer possible, and leaves a blank unambiguously meaning unanswered.

A respondent must be able to state the truth. Where the form asks which page a URL should be consolidated into, the answer is free text. Binding it to a drop-down of whatever rows happen to sit on the sheet would let a respondent nominate only from an arbitrary sample, and the answer that gets recorded would then be the closest available option rather than the correct one. Constrained input is right for a closed vocabulary and wrong for an open one; the same reasoning keeps `Webpage Custodian` unvalidated in §2.

The data region holds nothing but data. Drop-down source lists live on a dedicated `Lists` sheet, never in spare columns beside the records. Anything that reads a sheet's used range (Power Query, a copy-paste, an export) picks up those spare columns as though they were fields, and the resulting phantom columns are laborious to trace back. Separating them costs one sheet.

Every validation list contains every value the data contains. A list that omits a value still present in the records is worse than no list at all: it silently invalidates real rows and trains users to click through the warning. This is why the system labels of §2 appear on the `Lists` sheet as first-class entries, and it is enforced as rule 10 in Appendix A.

Every table in the model is bound to a visual. A table loaded into the model but referenced by nothing is dead weight that misleads whoever maintains the model next, because they have no way to tell an unused table from one whose visual was deleted by accident, so they either preserve it forever or remove something load-bearing. The performance table earns its place in the model by carrying the evidence behind `Very Low Page Views`.

7. Measures

Seven measures drive the dashboard cards, plus one calculated column.

On provenance. These are stated as specification (what each measure means and how it behaves under slicing) rather than extracted from the file.

Measure	Definition
Total Pages	Row count of the custodianship table. The denominator for both percentages.
URLs Discussed	Outcome rows where Date of Sunsetting Discussion is not blank.
Procedures Confirmed	Outcome rows where Sunsetting Procedure is not Yet To Be Confirmed .
% Procedures Confirmed	$\text{Procedures Confirmed} \div \text{Total Pages} .$
Custodian Count	Distinct count of Webpage Custodian .
Custodians Confirmed	Custodianship rows where Custodian Confirmed = Yes .
% Custodians Confirmed	$\text{Custodians Confirmed} \div \text{Total Pages} .$
Tracking Stage	Calculated column grouping Sunsetting Procedure into progress stages.

Three notes on reading them correctly; see Appendix B for the full semantics.

URLs Discussed and Procedures Confirmed are not the same quantity, and the gap between them is the most useful number on the page. A page can be discussed and then left at Another Discussion to be Held : discussed, but not confirmed. The gap is the programme's deferred backlog. Two measures are kept rather than one because process progress and decision progress move at different rates and a single "progress" measure would hide that.

Custodians Confirmed counts pages, not people. It sits next to Custodian Count , which counts accountable parties, and the adjacency invites a comparison that means nothing. % Custodians Confirmed is the honest reading: the share of the estate under confirmed ownership.

Custodian Count counts accountable parties, including the three system labels. That is consistent with how the column is defined in §2 and with every visual that groups by custodian; it is not a headcount and should not be quoted as one.

Appendix A: Integrity rules

These ten rules hold in the published data and should hold in any copy of this model. They are not assertions in prose: they are executable, and the build fails if any one of them breaks. Rules 6, 9 and 10 each earned their place by catching a genuine defect in the data generator during development, which is the only real evidence that a check does any work.

#	Rule
1	<code>URL Links</code> is unique within each of the three files.
2	Custodianship and outcomes contain an identical set of URLs.
3	Every performance URL exists in custodianship. The reverse does not hold.
4	No performance row exists for a URL flagged <code>NOT in Sitemap - Potential Sunset</code> .
5	<code>Custodian Confirmed = No</code> if and only if the custodian is <code>Yet to be Confirmed</code> .
6	<code>Webpage Custodian = Yet to be Confirmed</code> implies <code>Sunsetting Procedure = Yet To Be Confirmed</code> . An unowned page cannot carry a confirmed decision.
7	<code>Sunsetting Procedure = Yet To Be Confirmed</code> implies both <code>Reason</code> and <code>Date of Sunsetting Discussion</code> are blank.
8	<code>Reason</code> is blank for <code>Another Discussion to be Held and Digital Team Discretion</code> .
9	Every <code>Business Unit</code> maps to exactly one <code>Business Divisions</code> .
10	Every value in a controlled column appears on its workbook's <code>Lists</code> sheet.

Rule 6 is the one that matters most in use. It is the formal statement of the programme's central principle: **a page nobody owns cannot be marked for deletion**. Every other rule protects the model; this one protects the content.

Appendix B: Measure semantics and edge cases

Slicing. `Total Pages` counts stored-file handles alongside pages unless `Website Section` is sliced, so both percentages shift substantially the moment that slicer moves. This is correct behaviour (the estate genuinely includes both), but a figure quoted from an unsliced card and a figure quoted from a sliced one are not comparable, and neither is wrong.

Division by zero. `% Procedures Confirmed` and `% Custodians Confirmed` both divide by `Total Pages`. An empty slicer selection makes that divisor zero, so the divisor is guarded rather than left to return an error into a KPI card.

`Custodian Count` **under slicing** behaves as a distinct count, not a sum: filtering to two divisions that share a custodian returns fewer parties than adding the two division figures together. Standard for a distinct count, and routinely misread on a card that shows no context.

`Tracking Stage` is a calculated column rather than a measure because it groups rows for an axis: it assigns each of the ten procedures to a progress stage so the column chart has a meaningful order. Implementing it as a measure would leave the chart sorting stages alphabetically.

Appendix C: The published dataset in numbers

Every figure below is a property of the **synthetic** dataset. It is reproducible: the generator is seeded, and regenerating produces byte-identical output.

Estate. 12,710 URLs: Media Library 7,400 · Instruments 2,150 · Announcements 1,680 · Business 1,480. 48 custodian values (45 named individuals, 3 system labels), 39 business units (37 units, 2 system labels), 17 divisions (15 divisions, 2 system labels).

Decisions. To Be Decommissioned 4,237 · Another Discussion to be Held 2,765 · Yet To Be Confirmed 2,305 · To Be Migrated 1,628 · Digital Team Discretion 793 · Move To Client Portal 290 · Move To Group Site 275 · Temporarily Unpublished 265 · To Be Migrated (Content Update) 138 · Move To Clearing House Site 14.

Validation flags. In Sitemap 8,658 · NOT in Sitemap 2,641 · In Sitemap - New 1,411.

Field fill. 9,174 rows carry a discussion date · 4,956 a reason · 8,732 a comment.

Traffic. 9,727 URLs with records, 4,757,296 total views, median 12, maximum 1,600,000, 4,077 URLs under 10 views. Section medians: Business 86 · Instruments 18 · Announcements 12 · Media Library 8.

The 2,983 URLs with no traffic record are the 2,641 flagged `NOT in Sitemap` plus 342 pages published after the measurement window: the population §1 warns about.

Content Sunseting Programme Tracker · Data Dictionary. The schema, vocabularies and measures described here are those of a tracking system built and maintained by the author. All values shown are synthetic and describe a fictitious exchange.

Ndivhuwo · ndivhuwojse@gmail.com · [linkedin.com/in/ndivhuwo-makhavhu](https://www.linkedin.com/in/ndivhuwo-makhavhu) · github.com/MainDevWork